

# IMPORTING MODULES

```
In [1]: import numpy as np

# dataa split
from sklearn.model_selection import train_test_split

# model Evaluation
from sklearn import metrics

#navie bayesian and accuracy
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score

#pandas
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
plt.style.use('seaborn')
```

```
In [2]: import seaborn as sns
import plotly.express as px

# SMOTE
from imblearn.over_sampling import SMOTE

# scaling
from sklearn.preprocessing import StandardScaler

# tune
from sklearn.model_selection import RandomizedSearchCV
```

# READING THE DATASET

```
In [3]: data=pd.read_csv('D3.csv')
data
```

```
Out[3]:
```

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type
0	9046	Male	67	0	1	Yes	Private	Urban
1	31112	Male	80	0	1	Yes	Private	Rural
2	60182	Female	49	0	0	Yes	Private	Urban
3	1665	Female	79	1	0	Yes	Self-employed	Rural
4	56669	Male	81	0	0	Yes	Private	Urban
...	...	...	...	...	...	...	...	...
2538	5006	Female	46	0	0	Yes	Self-employed	Rural
2539	11250	Male	78	0	0	Yes	Self-employed	Rural
2540	41858	Female	63	0	1	Yes	Private	Rural

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type
2541	34965	Female	18	0	0	No	Private	Urban
2542	65748	Female	46	0	0	Yes	Private	Urban

2543 rows × 12 columns



## DATA ENCODING (PREPARATION)

```
In [4]: # convert string to numeric using map

# gender
data['gender'] = data['gender'].map({
    'Male': int(0),
    'Female': int(1),
    'Other': int(2)})

# ever_married
data['ever_married'] = data['ever_married'].map({
    'Yes': int(1),
    'No': int(0)})

# work_type
data['work_type'] = data['work_type'].map({
    'Private': int(3),
    'Self-employed': int(4),
    'Govt_job': int(2),
    'children': int(1),
    'Never_worked': int(0)})

# Residence_type
data['Residence_type'] = data['Residence_type'].map({
    'Urban': int(2),
    'Rural': int(1)})

# smoking_status
data['smoking_status'] = data['smoking_status'].map({
    'formerly smoked': int(1),
    'never smoked': int(2),
    'smokes': int(3),
    'Unknown': int(0)})
```

attributes used in the classification

```
In [5]: x=data[['gender','age','hypertension','heart_disease','ever_married','work_ty
y=data[['stroke']]
x=x.values
y=y.values
```

## SPLIT DATASET

```
In [6]: from sklearn.model_selection import train_test_split
from sklearn.datasets import load_iris
```

```
In [7]: X_train, X_test, y_train, y_test = train_test_split(x, y, test_size = 0.20, ra
```

```
In [8]: print(X_train.shape)
        print(X_test.shape)
```

```
(2034, 10)
(509, 10)
```

```
In [9]: y_train=y_train.flatten()
        y_test=y_test.flatten()
        print(y_train.shape)
        print(y_test.shape)
```

```
(2034,)
(509,)
```

## NAVIE'S BAYER

```
In [10]: from sklearn.naive_bayes import GaussianNB
```

```
In [11]: model = GaussianNB()
        model.fit(X_train,y_train)
        predictions=model.predict(X_test)
```

```
In [12]: print(np.unique(predictions))
```

```
[0 1]
```

```
In [13]: print('1. CONFUSION MATRIX\n',metrics.confusion_matrix(y_test, predictions))

        print("\n2. F1 SCORE")
        print('F1-score on Test set:\t',metrics.f1_score(y_test,predictions))

        print('\n3. OTHER METRICS')
        print(metrics.classification_report(y_test, predictions))

        # accuracy score
        train_score =model.score(X_train,y_train)
        test_score = model.score(X_test,y_test)

        print("\n4. TRAINING AND TEST ERRORS")
        print('Accuracy on Train set\t',train_score)
        print('Error on Train set\t',1-train_score)
        print('Accuracy on Test set\t',test_score)
        print('Error on Test set\t',1-test_score)
```

### 1. CONFUSION MATRIX

```
[[427  49]
 [ 23  10]]
```

## 2. F1 SCORE

F1-score on Test set: 0.21739130434782608

## 3. OTHER METRICS

	precision	recall	f1-score	support
0	0.95	0.90	0.92	476
1	0.17	0.30	0.22	33
accuracy			0.86	509
macro avg	0.56	0.60	0.57	509
weighted avg	0.90	0.86	0.88	509

## 4. TRAINING AND TEST ERRORS

Accuracy on Train set 0.8500491642084562  
 Error on Train set 0.14995083579154378  
 Accuracy on Test set 0.8585461689587426  
 Error on Test set 0.14145383104125742

# DECISION TREE

```
In [14]: from sklearn.tree import DecisionTreeClassifier
```

```
In [15]: model= DecisionTreeClassifier(random_state=42)
model.fit(X_train, y_train)
predictions = model.predict(X_test)
```

```
In [16]: print(np.unique(predictions))
```

```
[0 1]
```

```
In [17]: print('1. CONFUSION MATRIX\n',metrics.confusion_matrix(y_test, predictions))

print("\n2. F1 SCORE")
print('F1-score on Test set:\t',metrics.f1_score(y_test,predictions))

print('\n3. OTHER METRICS')
print(metrics.classification_report(y_test, predictions))

# accuracy score
train_score =model.score(X_train,y_train)
test_score = model.score(X_test,y_test)

print("\n4. TRAINING AND TEST ERRORS")
print('Accuracy on Train set\t',train_score)
print('Error on Train set\t',1-train_score)
print('Accuracy on Test set\t',test_score)
print('Error on Test set\t',1-test_score)
```

## 1. CONFUSION MATRIX

```
[[447  29]
 [ 29   4]]
```

## 2. F1 SCORE

F1-score on Test set: 0.12121212121212122

## 3. OTHER METRICS

	precision	recall	f1-score	support
0	0.94	0.94	0.94	476
1	0.12	0.12	0.12	33
accuracy			0.89	509
macro avg	0.53	0.53	0.53	509
weighted avg	0.89	0.89	0.89	509

## 4. TRAINING AND TEST ERRORS

Accuracy on Train set 1.0  
 Error on Train set 0.0  
 Accuracy on Test set 0.8860510805500982  
 Error on Test set 0.11394891944990182

## KNN KNeighborsClassifier

```
In [18]: from sklearn.neighbors import KNeighborsClassifier
```

Here we took k=4.

This model will use the four nearest neighbors to predict the value of a future data point.

```
In [19]: model = KNeighborsClassifier(n_neighbors = 4)
model.fit(X_train, y_train.ravel())
predictions = model.predict(X_test)
```

```
In [20]: print(np.unique(predictions))
```

```
[0 1]
```

```
In [21]: print('1. CONFUSION MATRIX\n',metrics.confusion_matrix(y_test, predictions))

print("\n2. F1 SCORE")
print('F1-score on Test set:\t',metrics.f1_score(y_test,predictions))

print('\n3. OTHER METRICS')
print(metrics.classification_report(y_test, predictions))

# accuracy score
train_score =model.score(X_train,y_train)
test_score = model.score(X_test,y_test)

print("\n4. TRAINING AND TEST ERRORS")
print('Accuracy on Train set\t',train_score)
print('Error on Train set\t',1-train_score)
print('Accuracy on Test set\t',test_score)
```

```
print('Error on Test set\t',1-test_score)
```

#### 1. CONFUSION MATRIX

```
[[473   3]
 [ 33   0]]
```

#### 2. F1 SCORE

F1-score on Test set: 0.0

#### 3. OTHER METRICS

	precision	recall	f1-score	support
0	0.93	0.99	0.96	476
1	0.00	0.00	0.00	33
accuracy			0.93	509
macro avg	0.47	0.50	0.48	509
weighted avg	0.87	0.93	0.90	509

#### 4. TRAINING AND TEST ERRORS

```
Accuracy on Train set 0.9321533923303835
Error on Train set    0.06784660766961648
Accuracy on Test set  0.9292730844793713
Error on Test set     0.07072691552062871
```

## ANN Artifical Neural Networks

```
In [22]: import tensorflow as tf
         from keras.models import Sequential
         from keras.layers import Dense
         import matplotlib.pyplot as plt
```

```
In [23]: # define keras model
         model=tf.keras.Sequential()

         model.add(tf.keras.layers.Dense(units=25,activation='relu'))

         model.add(tf.keras.layers.Dense(units=25,activation='relu'))

         model.add(tf.keras.layers.Dense(units=1,activation='sigmoid'))

         #compile keras model
         model.compile('adam','binary_crossentropy',metrics=['accuracy'])
```

## Training ANN Model

```
In [24]: #fitting ANN to training set
         model.fit(X_train,y_train,epochs=5)

         #accuracy
         accuracy=model.evaluate(X_train,y_train)
```

```
Epoch 1/5
64/64 [=====] - 0s 612us/step - loss: 4.2688 - accur
acy: 0.5320
```

```
Epoch 2/5
64/64 [=====] - 0s 717us/step - loss: 0.4407 - accur
acy: 0.9286
Epoch 3/5
64/64 [=====] - 0s 730us/step - loss: 0.2819 - accur
acy: 0.9264
Epoch 4/5
64/64 [=====] - 0s 685us/step - loss: 0.2545 - accur
acy: 0.9315
Epoch 5/5
64/64 [=====] - 0s 681us/step - loss: 0.2528 - accur
acy: 0.9288
64/64 [=====] - 0s 429us/step - loss: 0.2446 - accur
acy: 0.9267
```

In [25]:

```
#predictions
predictions = model.predict(X_test)
predictions = (predictions > 0.5)
```

In [26]:

```
print('1. CONFUSION MATRIX\n',metrics.confusion_matrix(y_test, predictions))

print("\n2. F1 SCORE")
print('F1-score on Test set:\t',metrics.f1_score(y_test,predictions))

print('\n3. OTHER METRICS')
print(metrics.classification_report(y_test, predictions))
print("\n 4.ACCURACY")
print(accuracy)
```

1. CONFUSION MATRIX

```
[[476  0]
 [ 33  0]]
```

2. F1 SCORE

F1-score on Test set: 0.0

3. OTHER METRICS

	precision	recall	f1-score	support
0	0.94	1.00	0.97	476
1	0.00	0.00	0.00	33
accuracy			0.94	509
macro avg	0.47	0.50	0.48	509
weighted avg	0.87	0.94	0.90	509

4.ACCURACY

[0.24461983144283295, 0.9267453551292419]

```
/home/pandu/my_project_dir/my_project_env/lib/python3.8/site-packages/sklearn
n/metrics/_classification.py:1248: UndefinedMetricWarning: Precision and F-sc
ore are ill-defined and being set to 0.0 in labels with no predicted samples.
Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
/home/pandu/my_project_dir/my_project_env/lib/python3.8/site-packages/sklearn
n/metrics/_classification.py:1248: UndefinedMetricWarning: Precision and F-sc
ore are ill-defined and being set to 0.0 in labels with no predicted samples.
Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
/home/pandu/my_project_dir/my_project_env/lib/python3.8/site-packages/sklearn
n/metrics/_classification.py:1248: UndefinedMetricWarning: Precision and F-sc
ore are ill-defined and being set to 0.0 in labels with no predicted samples.
Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```